

Introduction to Parallel Distributed Processing models

Deon T. Benton

003-Hebb and Delta rules

A note about **activation functions**

- We briefly discussed a variety (threshold, linear, sigmoid, tanh, etc) but zoomed in on sigmoid
- Often to simplify the math it is useful to think about threshold functions where activations are -1 if input is negative and +1 if input is positive
- Can view this as indicating whether a unit is *more* or *less* active than its tonic firing rate.
- Other times it is mathematically useful to think about *linear* activation functions (activation = net input).
- These have less clear neural interpretations but are a useful stepping-stone for understanding learning algorithms.

A long-standing view of learning/memory

- **William James, 1890:** Memory arises from association and generalization
 - *Association:* Two things that regularly occur together (“contiguity”) become associated in memory; encountering one brings the other to mind.
 - e.g.: image of German Shepard + word “dog”
 - *Generalization:* Same association will be promoted by new stimuli when they are “similar” to one of the associates
 - e.g.: image of a Labrador evokes word “dog” b/c Labrador is similar to German Shepard
- How might such learning occur in a neural network?

Associative learning

- A given item (word, image, sound, motor action, feeling, etc) is represented as a pattern of activity over units
- Two items (e.g. word, picture) represented by two patterns over different units, connected via weights
- When they are “contiguous,” both patterns simultaneously active
- Weights must change so that one pattern presented alone generates the other pattern

A mechanism: Hebbian learning

- Hebb:
 - When an axon of cell A is near enough to excite a cell B and *repeatedly and consistently takes part in firing it*, some growth process or metabolic change takes place in one or both cells such that A's efficacy, as one of the cells firing B, is increased.
- Minimal Hebb rule:
 - When there is a synapse between cell A and cell B, increment the strength of the synapse whenever A and B fire together (or in close succession).
- In math:
 - $\Delta w_{ij} = \varepsilon a_i a_j$ (outer product)
 - ...where w_{ij} is the weight from i to j, ε is a constant and a_i and a_j are the activations of the connected units

A note on *superpositional weight changes*

- Changes to the weight matrix are superpositional
 - changes made to a weight matrix, w , by the outer product of the activity between an input and output pattern pair at time t will be superimposed on the weight changes caused by another pair at time $t + 1$.

Thus,

$$W = \begin{pmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{pmatrix}$$

$$(u_1 y_1) \rightarrow \begin{pmatrix} 1 & 1 & 0 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{pmatrix}$$

$$(u_2 y_2) \rightarrow \begin{pmatrix} 1 & 1 & 0 \\ 1 & 1 & 1 \\ 1 & 1 & 0 \end{pmatrix}$$

$$\begin{pmatrix} 2 & 2 & 0 \\ 1 & 2 & 1 \\ 2 & 1 & 0 \end{pmatrix}$$



A quick aside: The **inner (or dot)** product *between two vectors*

- If we have two vectors, u and v , where

$$u = \begin{pmatrix} u_1 \\ u_2 \\ u_3 \end{pmatrix} \text{ and } v = \begin{pmatrix} v_1 \\ v_2 \\ v_3 \end{pmatrix}$$

Then their *inner* product is

$$u^T v = (u_1 \ u_2 \ u_3) \begin{pmatrix} v_1 \\ v_2 \\ v_3 \end{pmatrix} = u_1 v_1 + u_2 v_2 + u_3 v_3$$

A quick aside: The **outer product** *between two vectors*

- If we have two vectors, u and v , where

$$u = \begin{pmatrix} u_1 \\ u_2 \\ u_3 \end{pmatrix} \text{ and } v = \begin{pmatrix} v_1 \\ v_2 \\ v_3 \end{pmatrix}$$

Then their *outer* product is

$$uv^T = \begin{pmatrix} u_1 \\ u_2 \\ u_3 \end{pmatrix} (v_1 \ v_2 \ v_3) = \begin{matrix} u_1 v_1 & u_1 v_2 & u_1 v_3 \\ u_2 v_1 & u_2 v_2 & u_2 v_3 \\ u_3 v_1 & u_3 v_2 & u_3 v_3 \end{matrix}$$

If $w_{ij} = 0$ initially, after a set of n training trials on patterns (indexed by p) where $\Delta w_{ij} = \epsilon a_i a_j$,

$$w_{ij} = \epsilon \sum_{p=1}^n a_i^{[p]} a_j^{[p]}$$

notation: $a_i^{[p]}$ is unit i 's activation in pattern p

Suppose a_i and a_j take on values of $+1$ or -1

- If a_i and a_j are *perfectly correlated* (always the same), $a_i^{[p]} a_j^{[p]} = 1$, so

$$w_{ij} = \epsilon n$$

- If a_i and a_j are *perfectly anticorrelated* (always differ), $a_i^{[p]} a_j^{[p]} = -1$, so

$$w_{ij} = -\epsilon n$$

- If a_i and a_j are *uncorrelated* (differ as often as same)

$$w_{ij} = \epsilon \left(\frac{n}{2}(+1) + \frac{n}{2}(-1) \right) = 0$$

- If a_i and a_j are *partially correlated* (e.g., 3/4 same and 1/4 different)

$$w_{ij} = \epsilon n \left(\frac{3}{4}(+1) + \frac{1}{4}(-1) \right) = \frac{1}{2} \epsilon n$$

- Thus $w_{ij} \propto \text{correlation}(a_i, a_j)$

Statistical correlation

- $$r_{xy} = \frac{\sum_d (x_d - \bar{x})(y_d - \bar{y})}{\sqrt{(\sum_d (x_d - \bar{x})^2)(\sum_d (y_d - \bar{y})^2)}}$$
- Denominator just normalizes to $[-1,1]$ range, so:
- $r_{xy} \propto \sum_d (x_d - \bar{x})(y_d - \bar{y})$
- In fact if activations are in $[-1,1]$ and each unit has a mean activation of 0 across all patterns, then:
 - $a_i = (x_d - \bar{x}) \rightarrow a_i = (x_d - 0) = a_i = x_d$
 - $a_j = (y_d - \bar{y}) \rightarrow a_j = (y_d - 0) = a_j = y_d$
 - $w_{ij} \propto r_{xy} \rightarrow a_i a_j \propto x_d y_d$

If test pattern p' is orthogonal to all training patterns p , $dp(p', p) = 0$ for all p , so

$$a_j^{[p']} = \epsilon \sum_p a_j^{[p]} dp(p', p) = \epsilon \sum_p a_j^{[p]} 0 = 0$$

If all training patterns are orthogonal to each other (and assuming $\epsilon = 1$), then

- If p' is one of the training patterns (say p^*), recall is perfect:

$$a_j^{[p']} = a_j^{[p^*]} dp(p^*, p^*) + \sum_{p \neq p^*} a_j^{[p]} dp(p', p) = a_j^{[p^*]} + \sum_{p \neq p^*} a_j^{[p]} 0 = a_j^{[p^*]}$$

- If p' is *similar* to only one training pattern (p^*) and orthogonal to the rest, the output is $a_j^{[p^*]}$ scaled by the degree of similarity:

$$a_j^{[p']} = a_j^{[p^*]} dp(p', p^*) + \sum_{p \neq p^*} a_j^{[p]} dp(p', p) = a_j^{[p^*]} dp(p', p^*)$$


In general, the output to any input pattern is a weighted combination of the outputs of all trained patterns, scaled by their similarity to the input.

- If the combination agrees with $a_j^{[p']}$, this is **facilitation** (or generalization if p' is novel)
- If the combination disagrees with $a_j^{[p']}$, this is **interference** (or poor generalization)

Houston, we have a problem!

- Unit-wise correlations are often insufficient to produce the correct output response

A




B

Input		Output
0	1	2
+	+	+
+	-	+
-	+	-
-	-	-

Final weights: +4, 0

Houston, we have a problem!

- Unit-wise correlations are often insufficient to produce the correct output response




D

	Input				Output
	0	1	2	3	4
	+	-	+	-	+
	+	+	+	+	+
	+	+	+	-	-
	+	-	-	+	-

Final weights: +4, 0, +3, 0

Houston, we have a problem!

- Unit-wise correlations are often insufficient to produce the correct output response



Error-correcting learning: Delta rule

Change weights so as to reduce difference between actual output (a_j) and **target** output (denoted t_j)

$$\Delta w_{ij} = \epsilon (t_j - a_j) a_i$$

- “Delta”: difference between output and target
 - Also called Widrow-Hoff rule, LMS (least mean squared)
 - Related to perceptron convergence procedure (Rosenblatt)
- Similar to correlation with *error*
- Hebb rule: $\Delta w_{ij} = \epsilon t_j a_i$ (where t_j is activation “clamped” on the output unit)

Learning on orthogonal patterns (one pass): Delta = Hebb

Delta rule: $\Delta w_{ij} = \epsilon (t_j - a_j) a_i$ (assume linear units: $a_j = n_j$)

Note: Delta = Hebb if $a_j = 0$

For first pattern p_1 , $w_{ij} = 0$ so $a_j^{[p_1]} = n_j^{[p_1]} = 0$, and

$$\Delta w_{ij} (= w_{ij}) = \epsilon (t_j^{[p_1]} - 0) = t_j^{[p_1]} a_i^{[p_1]}$$

Hebb rule with target as output activation

For p_2 , $a_j^{[p_2]} = \sum_i a_i^{[p_2]} w_{ij} = \sum_i a_i^{[p_2]} (t_j^{[p_1]} a_i^{[p_1]}) = t_j^{[p_1]} \sum_i a_i^{[p_2]} a_i^{[p_1]} \underline{\sum_i a_i^{[p_2]} a_i^{[p_1]}}$ (dot product of p_1 and p_2)

Since p_1 and p_2 are orthogonal, $\sum_i a_i^{[p_2]} a_i^{[p_1]} = 0$, so $a_j^{[p_2]} = 0$. Thus

$$\Delta w_{ij} = t_j^{[p_2]} a_i^{[p_2]} \quad w_{ij} = t_j^{[p_1]} a_i^{[p_1]} + t_j^{[p_2]} a_i^{[p_2]}$$

Hebb rule again

In fact, $a_j^{[p]} = 0$ for the first presentation of each training pattern p , so at the end of one sweep through all the patterns:

$$w_{ij} = \epsilon \sum_p (t_j^{[p]} - a_j^{[p]}) a_i^{[p]} = \epsilon \sum_p t_j^{[p]} a_i^{[p]}$$

This is just **Hebbian learning** using targets t_j as output activations (a_j).

Note that the Delta rule is inherently *multi-pass* ($a_j \neq 0$ on subsequent presentations)

- Weight changes caused by one pattern affect error on others

Effects of training on response to input patterns

Calculated in terms of *changes* to activations for pattern p' caused by training on single pattern p :

$$\begin{aligned}\Delta a_j^{[p']} &= \sum_i a_i^{[p']} \Delta w_{ij} \\ &= \sum_i a_i^{[p']} \epsilon \left(t_j^{[p]} - a_j^{[p]} \right) a_i^{[p]} \\ &= \epsilon \left(t_j^{[p]} - a_j^{[p]} \right) \sum_i a_i^{[p']} a_i^{[p]} \\ &= \epsilon \left(t_j^{[p]} - a_j^{[p]} \right) dp(p', p)\end{aligned}$$

- If p and p' are orthogonal, training on p will have no effect on p'
- If p and p' are not orthogonal, training on p will affect performance on p' (weighted by similarity) which may be good (generalization) or bad (interference)

When does the Delta rule succeed or fail?

Delta rule is *optimal*

- Will find a set of weights that produces zero error if such a set exists

Need to distinguish “succeed” = zero error from “succeed” = correct binary classification

Guaranteed to succeed (zero error) if input patterns are **linearly independent** (LI)


- No pattern can be created by recombining scaled versions of the others
(i.e., there is *something unique* about each pattern; cf. Hebb: no similarity)
- Orthogonal patterns are linearly independent (LI is a weaker constraint)
- Linearly independent patterns can be similar as long as other aspects are unique

Succeed at binary classification of outputs: **Linear separability**

Linear separability

Delta rule is guaranteed to succeed at binary classification if the task is **linearly separable**

- Weights define a plane (line for two input units) through input (state) space for which $n_j = 0$
- Must be possible to position this plane such that all patterns requiring $n_j < 0$ are on one side and all patterns requiring $n_j > 0$ are on the other side
- Property of the relationship between input and target patterns
- **AND** and **OR** are linearly separable but **XOR** is not




$$n_j = a_1 w_1 + a_2 w_2 + b_j = 0$$


$$a_2 = -\frac{w_1}{w_2} a_1 - \frac{b_j}{w_2}$$

$$(y = a \ x + b)$$

XOR



AND




$$n_j = a_1 w_1 + a_2 w_2 + b_j = 0$$


$$a_2 = -\frac{w_1}{w_2} a_1 - \frac{b_j}{w_2}$$

$$(y = a x + b)$$

OR



XOR




$$n_j = a_1 w_1 + a_2 w_2 + b_j = 0$$


$$a_2 = -\frac{w_1}{w_2} a_1 - \frac{b_j}{w_2}$$

$$(y = a x + b)$$


AND



OR



XOR





XOR with extra dimension





XOR task can be converted to one that is linearly separable by adding a new “input”

- Corresponds to a third dimension in state space
- Task is no longer XOR

Inputs	Output
0 0 0	0
0 1 0	1
1 0 0	1
1 1 1	0



XOR with intermediate (“hidden”) units



- Intermediate units can re-represent input patterns as new patterns with **altered similarities**
- Targets which are not linearly separable in the input space can be linearly separable in the intermediate representational space
- Intermediate units are called “hidden” because their activations are not determined directly by the training environment (inputs and targets)